

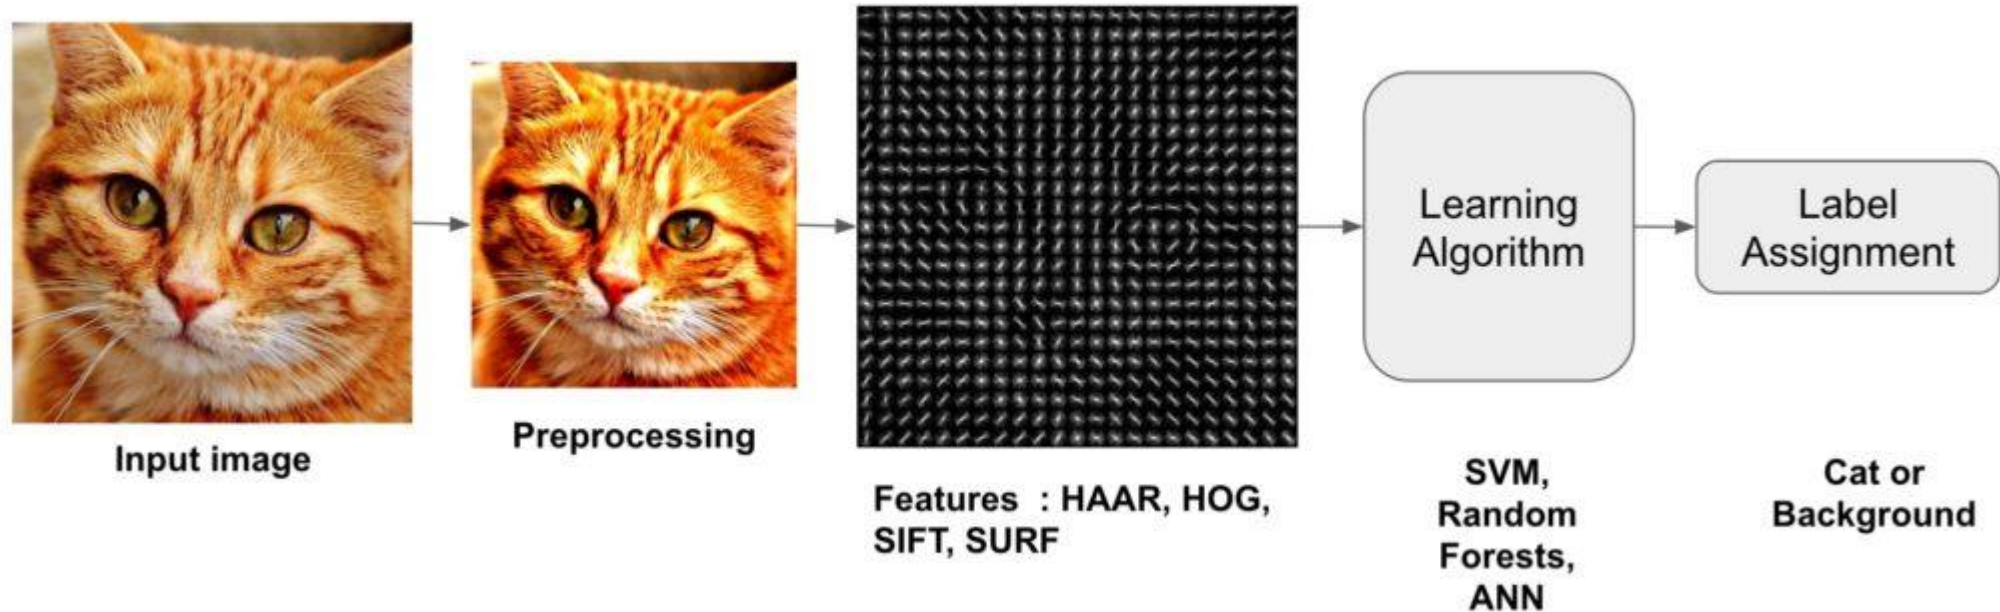
CNN Image Classification

PIRL, POSTECH

Hanul Roh

Recap: Traditional Image Classification


- A core task in Computer Vision



Data -> Feature extraction -> Classification

Recap: Traditional Image Classification


- A core task in Computer Vision



Data -> Feature extraction -> Classification

Convolution Neural Network

- The building of an **integrated feature extraction and classification**
- **Learning to extract features** from given training datasets
 - Receptive field (Locality)
 - Parameter sharing
 - Convolution + Subsampling -> FC layers -> Output



ConvNet Architecture

Convolution Layer

Compute the output of neurons that are connected to local regions

Activation Layer

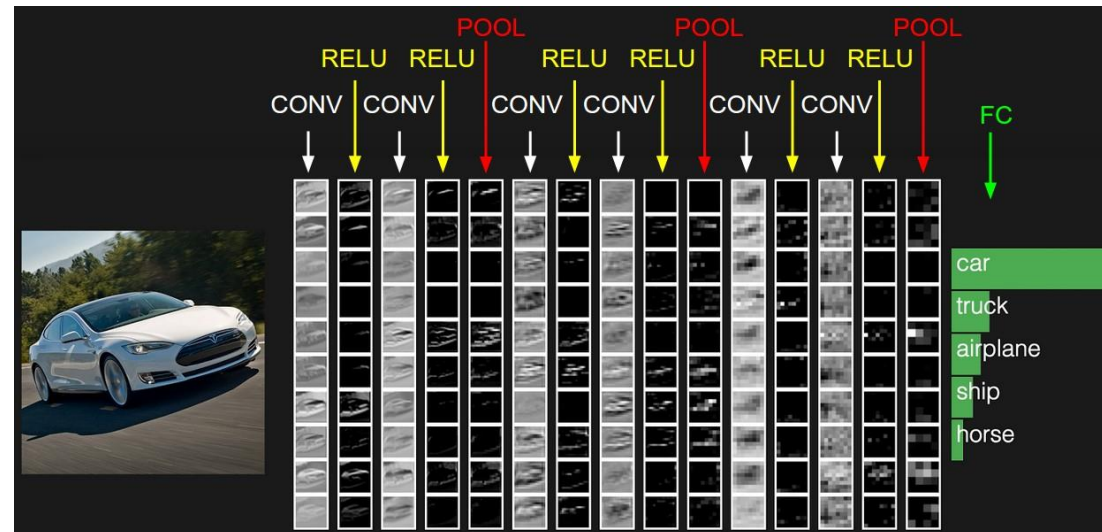
*Elementwise activation function
 $\max(0, f(x))$: nonlinearity*

Pooling Layer



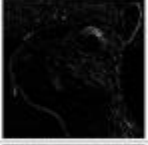




*Downsampling operation
2x2 input -> maximum value*

Fully-connected Layer

*Compute class scores resulting in
volume size $[1 \times 1 \times N]$*



Convolution Operation

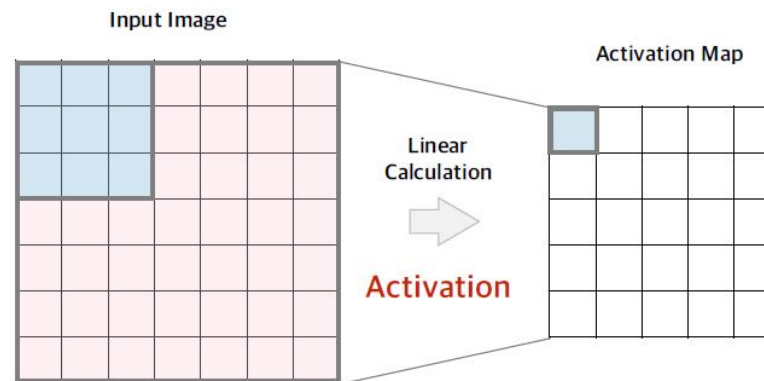

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

Convolution Layer

- Convolution layer demo here:

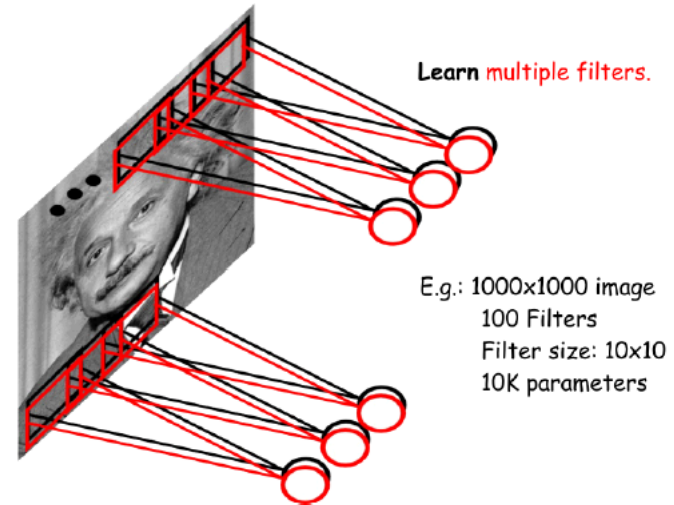



<http://cs231n.github.io/convolutional-networks/#layers>

<Linear Calculation of Filters on Image>



(Recall) In this example, only 9 value of weights are used to make activation map! → **Weight Sharing**

Convolution Layer



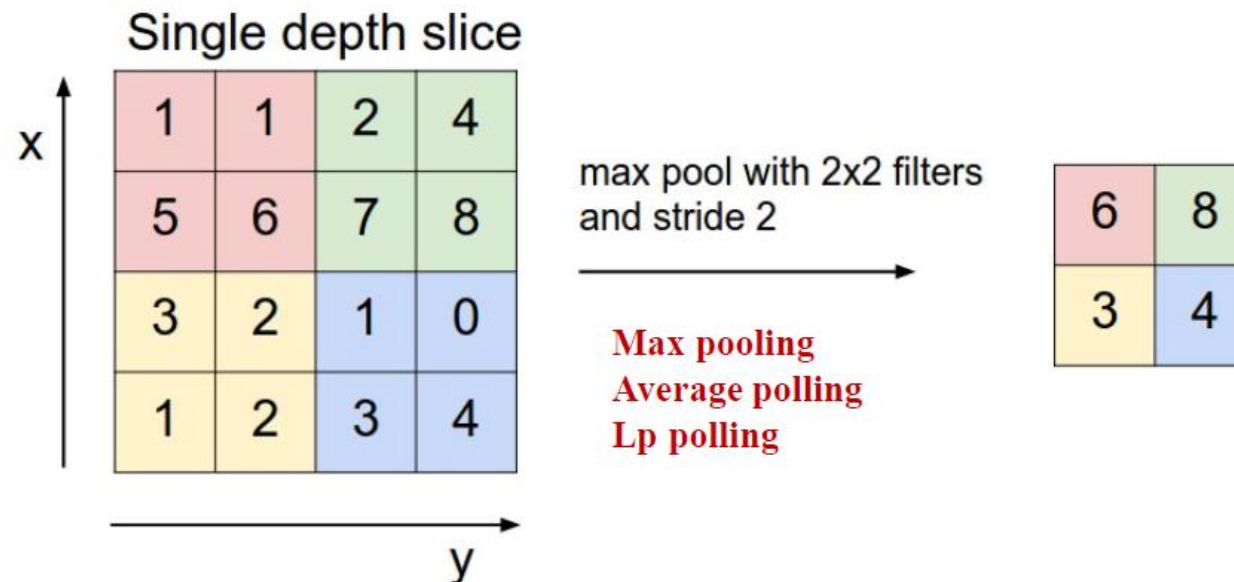
Convolutional layer

Convolution Layer

- Local Connection
 - Restricting the network architecture through the use of local connections known as receptive fields
- Parameter Sharing
 - All neurons in a feature share the same weights
 - Reduce the number of free parameters in the input images


Pooling Layer

- Reduce previous feature map size from $[2 \times 2]$ -> 1 value
 - Mean pooling, **Max pooling**, and so on
- Make the representations smaller and more manageable
- Operates over each activation map **independently**



Pooling Layer

- Reduce previous feature map size from $[2 \times 2] \rightarrow 1$ value
 - Mean pooling, **Max pooling**, and so on
- Make the representations smaller and more manageable
- Operates over each activation map **independently**




Pooling Layer

Subsampling

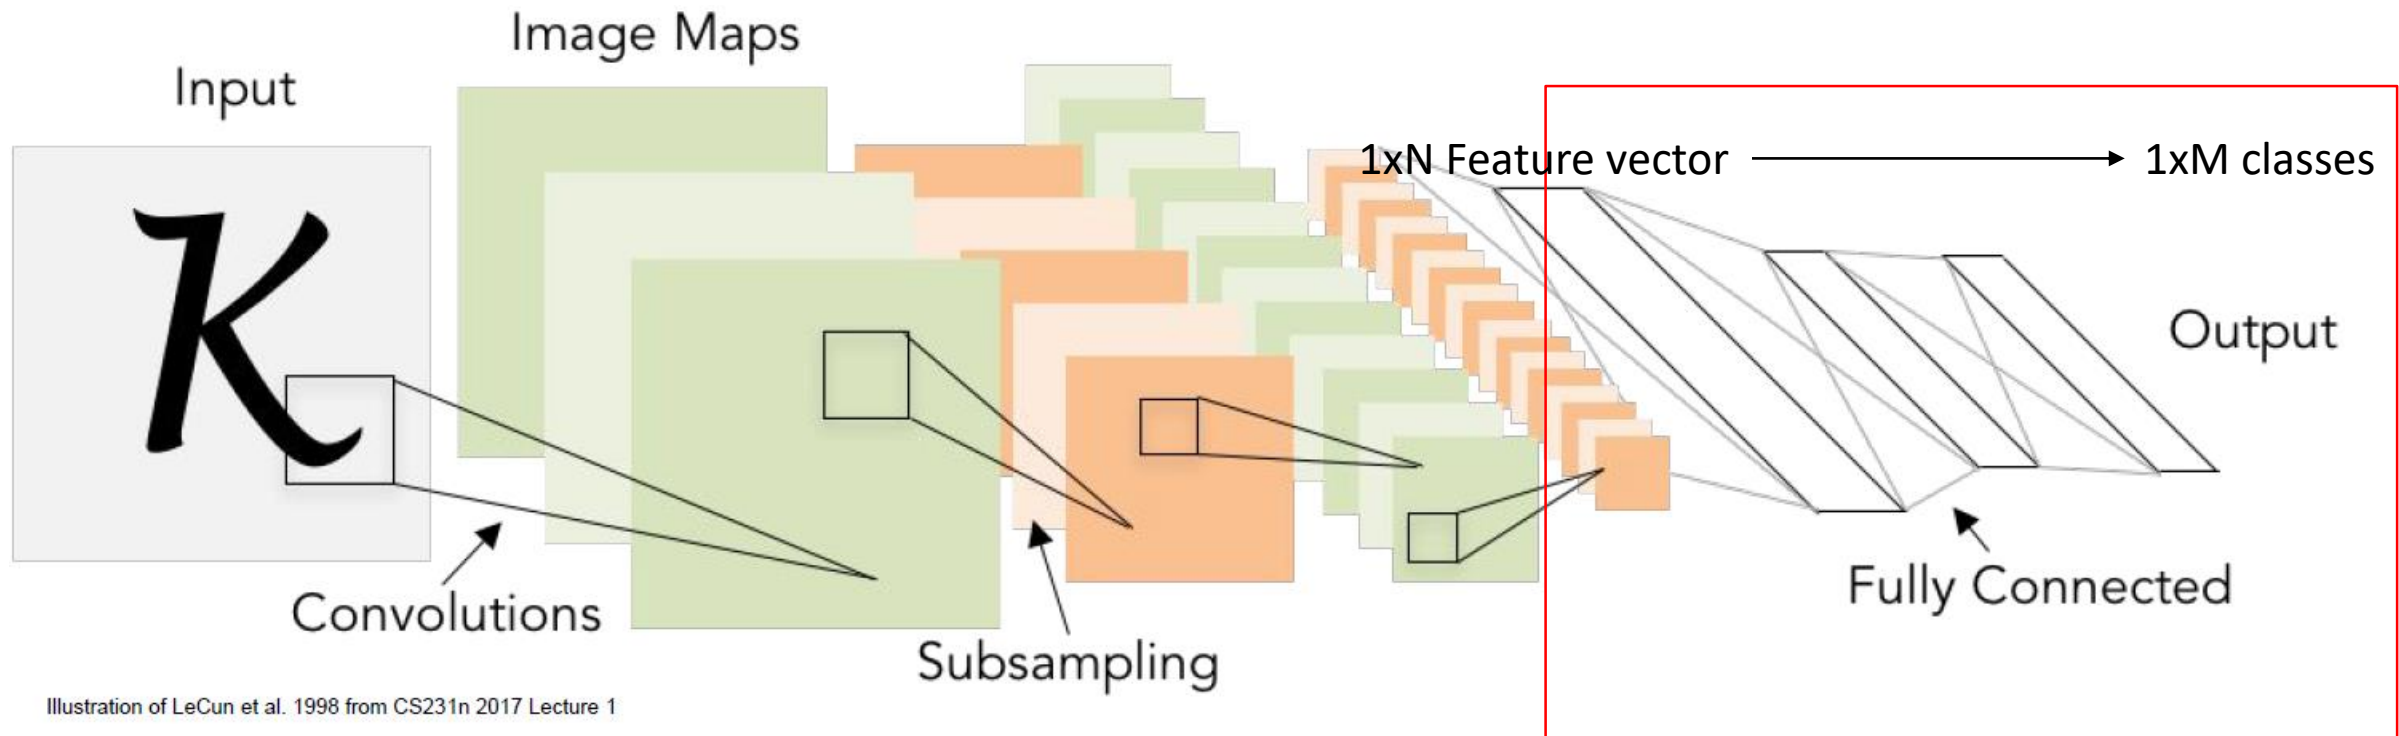
- Overlapping vs. non-overlapping
- Max pooling vs. average pooling

Translation invariance

- Pooling을 하면 작은 변화에는 둔감해짐
- 따라서, 이에 강인한 특징을 생성할 수 있음





Fully-connected Layer

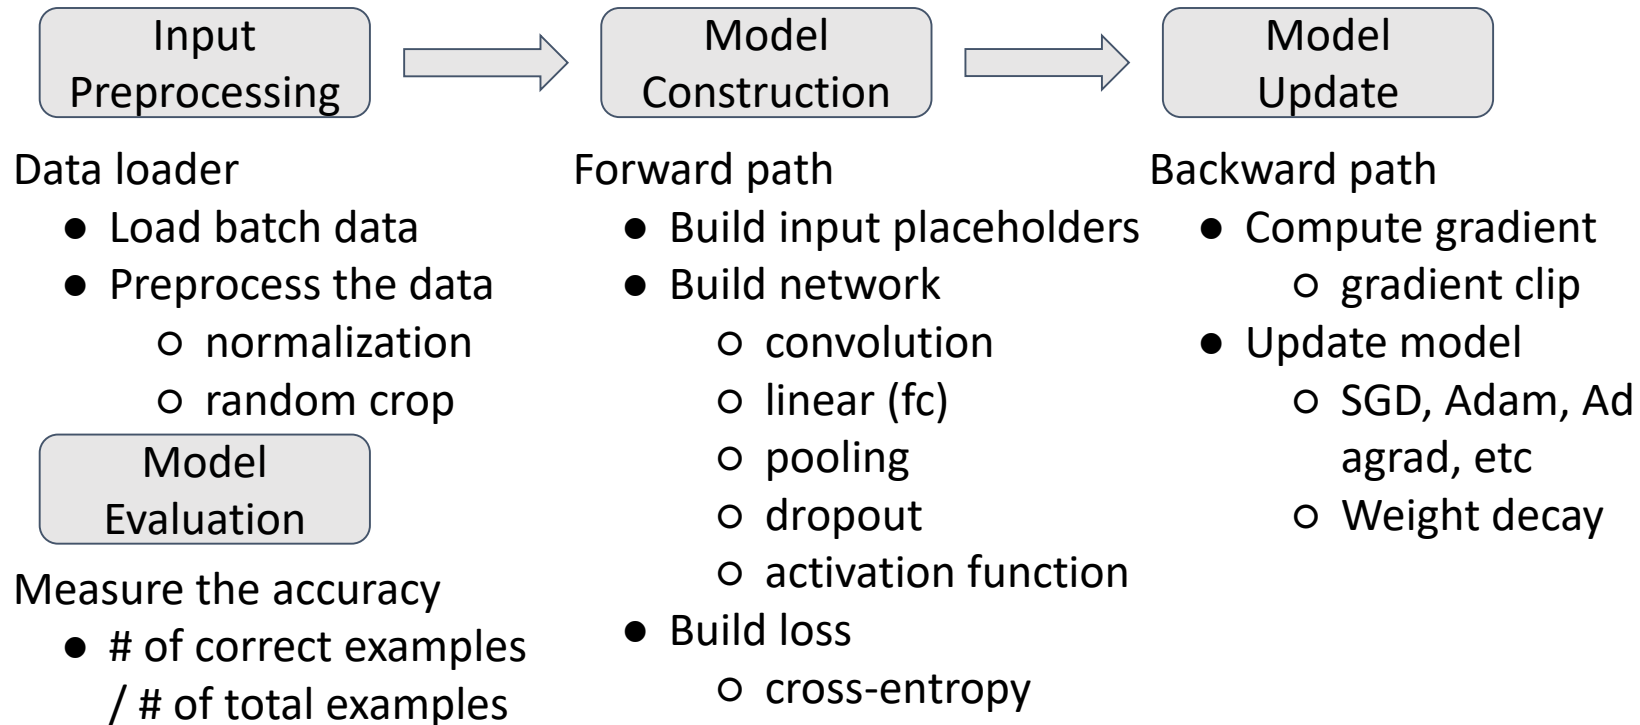


Activation Function

- Gradient vanishing problem -> ReLu!



CNN Image Classification



Input Preprocessing

- CIFAR-10 Dataset
 - **10 classes** - airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
 - **32x32 color images** with 6,000 images per class (50,000 training / 10,000 test)
 - **1x3072 vector** shape per each image

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck

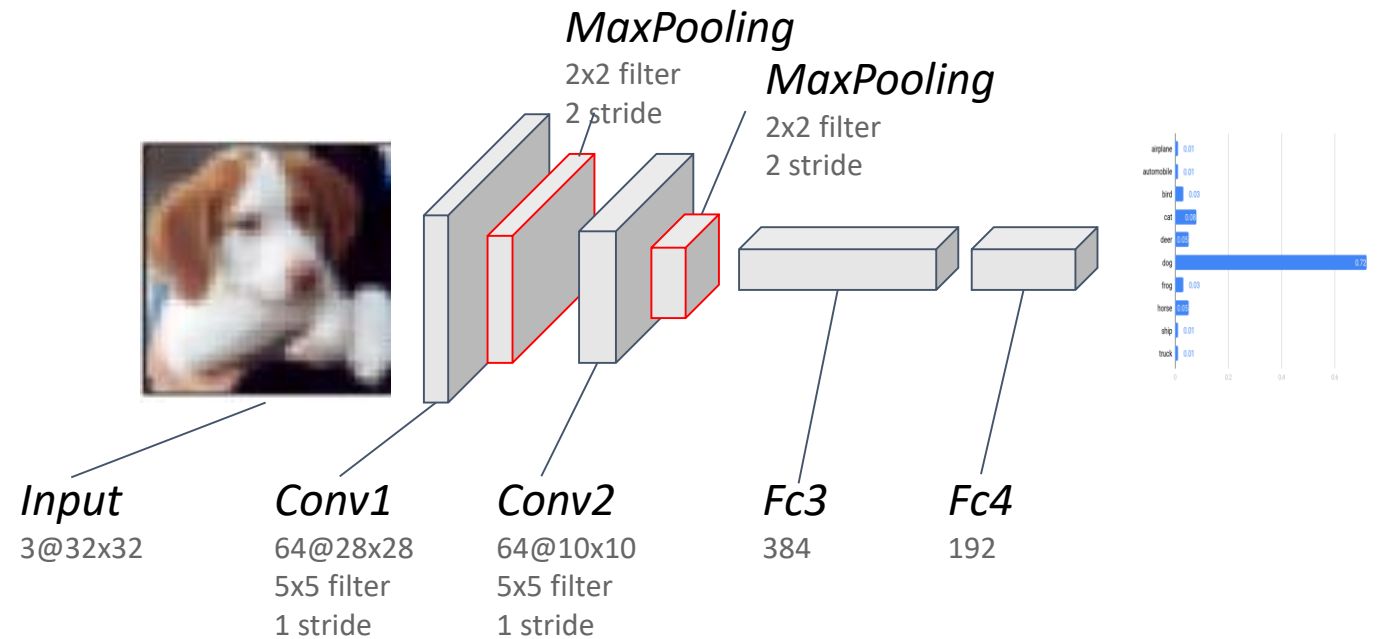


Input Preprocessing

- Data loader (cifar10_loader.py) does
 - Reshape a given vector of 3,072 size into a tensor of 3x32x32 size (RGB channels, width, height)
 - Transpose (channel, width, height) into (width, height, channel)
 - Load data of batch size using `get_batch()`

Model Construction

- In this practice,
 - 2 convolution layers + 2 fully connected layers + classification layer



Model Construction

- How to implement the layers?

Convolutional Layer - `tf.nn.conv2d()`

Create variables

```
filter = tf.get_variable("weights", [filter_size, filter_size, inp_dim, out_dim])  
bias = tf.get_variable("biases", [out_dim])
```

Compute convolution

```
conv = tf.nn.conv2d(input, filter)  
conv = tf.nn.bias_add(conv, bias)
```

Activation function (relu, tanh)?

```
conv = tf.nn.relu(conv)  
conv = tf.tanh(conv)
```

Max pooling?

kernel size 2 and stride 2

```
pool = tf.nn.max_pool(conv, ksize=[1,2,2,1], strides=[1,2,2,1])
```

Fully Connected Layer - `tf.matmul()`

Create variables

```
weight = tf.get_variable("weights", [inp_dim, out_dim])  
bias = tf.get_variable("biases", [out_dim])
```

Compute convolution

```
fc = tf.matmul(input, weight)  
fc = tf.nn.bias_add(fc, bias)
```

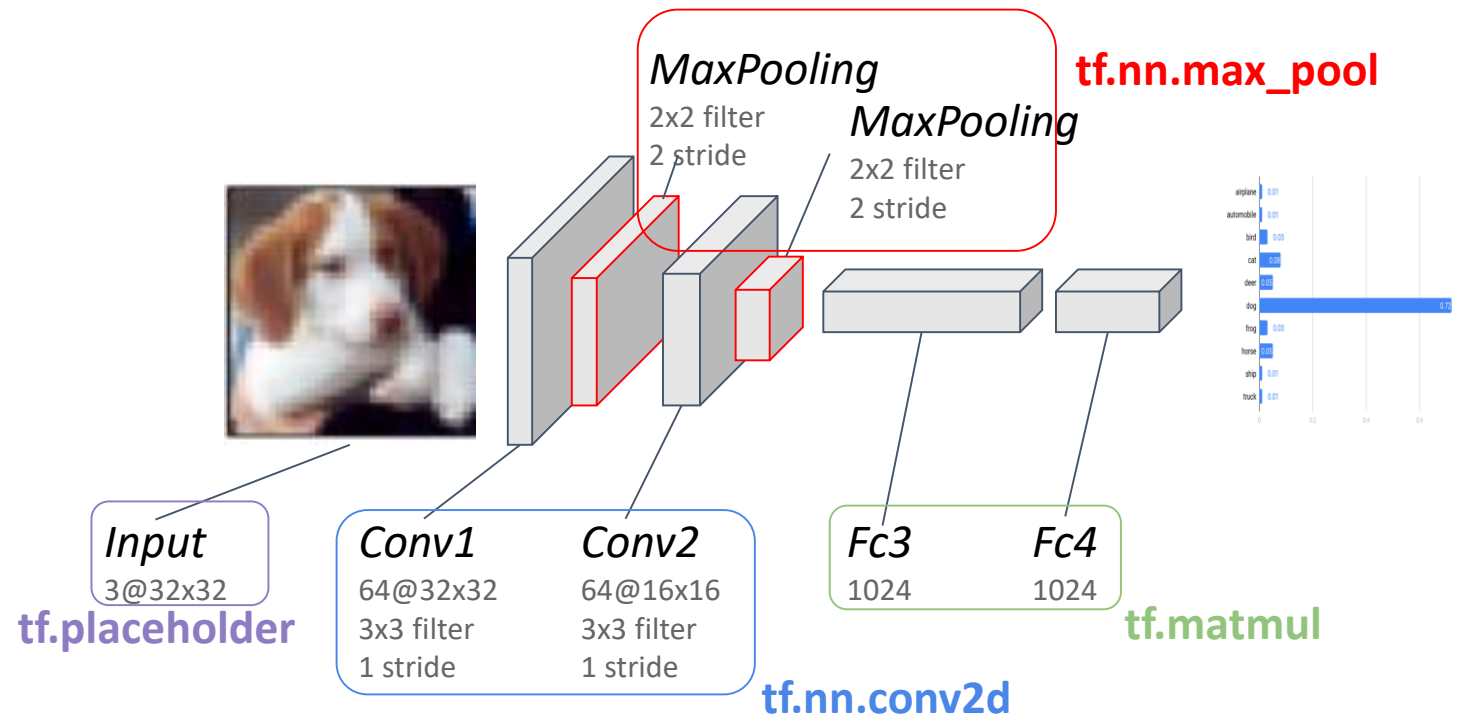
```
fc = tf.nn.relu(fc)  
fc = tf.tanh(fc)
```

Dropout?

```
drop_fc = tf.nn.dropout(fc, keep_prob)
```

Model Construction

- How to implement the layers in TensorFlow?



Model Construction

- We should always define variables manually? **No**
- There are three libraries based on TensorFlow which are **simple and concise!**

```
# Create variables
filter = tf.get_variable("weights", [filer_size,filter_size,
inp_dim, out_dim])
bias = tf.get_variable("biases", [out_dim])

# Compute convolution
conv = tf.nn.conv2d(input, filter)
conv = tf.nn.bias_add(conv, bias)

# Apply activation function
conv = tf.nn.relu(conv)
conv = tf.tanh(conv)

# Apply dropout
drop_conv = tf.nn.dropout(conv, keep_prob)
```


→
wrapper function

1. TFLearn
 - `tf.learn.layers.conv.conv_2d()`
2. Keras
 - `tf.contrib.keras.layers.Conv2D()`
3. TF-Slim
 - `tf.contrib.slim.conv2d()`

Note that we do not cover above libraries in this course

Model Update - Learning

- Define loss function (cross-entropy loss)
- Define optimizer (Adam)



Model Update - Learning

- How to define loss function and optimizer in TensorFlow?

Loss Function

Compute cross-entropy loss

```
cross_entropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels, logits)
loss = tf.reduce_mean(cross_entropy)
```

- cross-entropy with softmax

- Softmax:

$$S(y_i) = \frac{e^{y_i}}{\sum_{k=1} e^{y_k}}$$

cross-entropy: $H_y(y) = - \sum_i y_i \log(\text{softmax}(y_i))$

Now, we can train model with following line

```
sess = tf.Session()
sess.run(fetches=train_op, feed_dict=inputs)
```

Optimizer

Create optimizer

```
optimizer = tf.train.AdamOptimizer(learning_rate)
```

Two methods exists

1. compute gradients and update model

```
grads_vars = optimizer.compute_gradients(loss)
train_op = optimizer.apply_gradients(grads_vars)
```

2. Update model with one function

```
train_op = optimizer.minimize(loss)
```

* other optimizers?

- `tf.train.GradientDescentOptimizer()`
- `tf.train.RMSPropOptimizer()`
- `radOptimizer()`

Misc

How to show the variables defined in the network (graph)?

```
for var in tf.global_variables():  
    print(var)
```

* global_variables() returns the variables that are created by tf.Variable() or tf.get_variable()

* We can show operations using tf.get_default_graph().get_operations()

```
Tensor("conv1/weights/read:0", shape=(5, 5, 3, 64), dtype=float32)  
Tensor("conv1/biases/read:0", shape=(64,), dtype=float32)  
Tensor("conv2/weights/read:0", shape=(5, 5, 64, 64), dtype=float32)  
Tensor("conv2/biases/read:0", shape=(64,), dtype=float32)  
Tensor("fc3/weights/read:0", shape=(1600, 384), dtype=float32)  
Tensor("fc3/biases/read:0", shape=(384,), dtype=float32)  
Tensor("fc4/weights/read:0", shape=(384, 192), dtype=float32)  
Tensor("fc4/biases/read:0", shape=(192,), dtype=float32)  
Tensor("fc5/weights/read:0", shape=(192, 10), dtype=float32)  
Tensor("fc5/biases/read:0", shape=(10,), dtype=float32)
```

How to access weight parameters?

```
sess = tf.Session()  
g = tf.get_default_graph()  
w_tensor = g.get_tensor_by_name("fc5/weights/read:0")  
w = w_tensor.eval(session=sess)  
print(w.shape)  
print(w[0, :10])
```

```
# create session  
# get the current graph model  
# obtain a tensor variable using the name  
# get the tensor by running session
```

```
(192, 10)  
[ 0.00417436  0.01035099  0.00198316  0.00643856  0.00888029  0.00903297  
 -0.0122347  -0.01635222 -0.01234471  0.00342033]  
[ 0.00417436  0.01035099  0.00198316  0.00643856  0.00888029  0.00903297  
 -0.0122347  -0.01635222 -0.01234471  0.00342033]
```

Check the code `image_classification.ipynb`

- Train the model and evaluate it

Exercises

- Implement simple CNN architecture and learn CIFAR-10 datasets

SimpleCNN
9 weight layers
Input: CIFAR-10 32x32 RGB image
Conv3x3-64 Conv3x3-64
Maxpool
Conv3x3-128 Conv3x3-128 Conv3x3-128
Maxpool
FC-1024
FC-1024
FC-10
Soft-max