

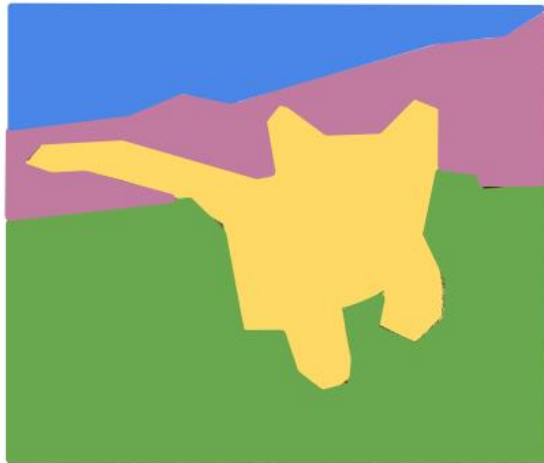
Object Detection

PIRL, POSTECH

Hanul Roh

Computer Vision Tasks

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

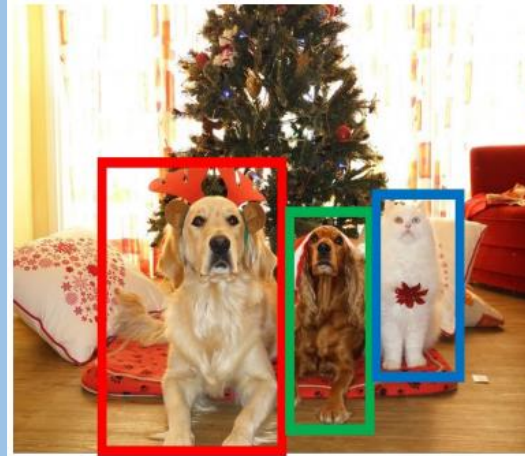
Classification + Localization



CAT

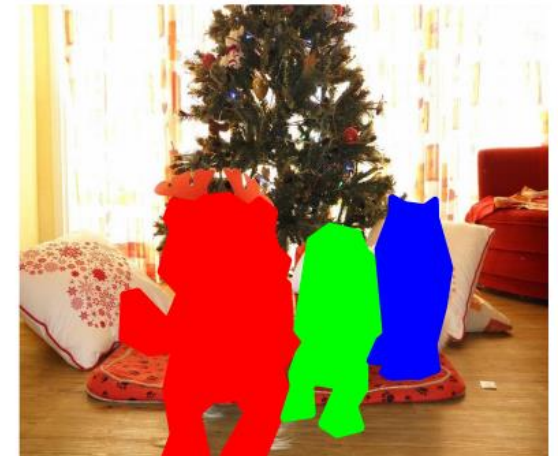
Single Object

Object Detection



DOG, DOG, CAT

Instance Segmentation



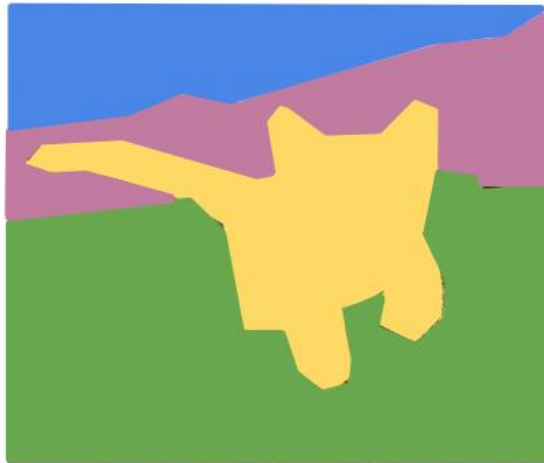
DOG, DOG, CAT

Multiple Object

[This image is CC0 public domain](#)

Computer Vision Tasks

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

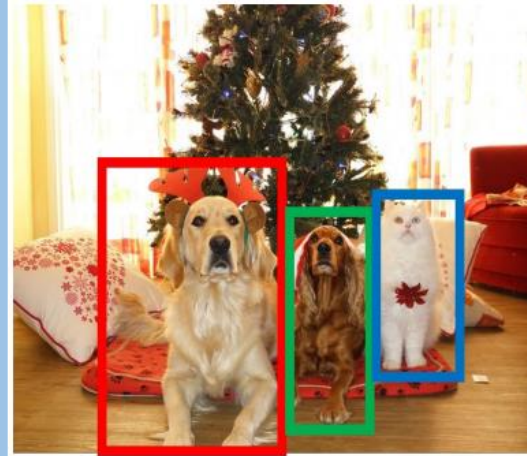
Classification + Localization



CAT

Single Object

Object Detection



DOG, DOG, CAT

Multiple Object

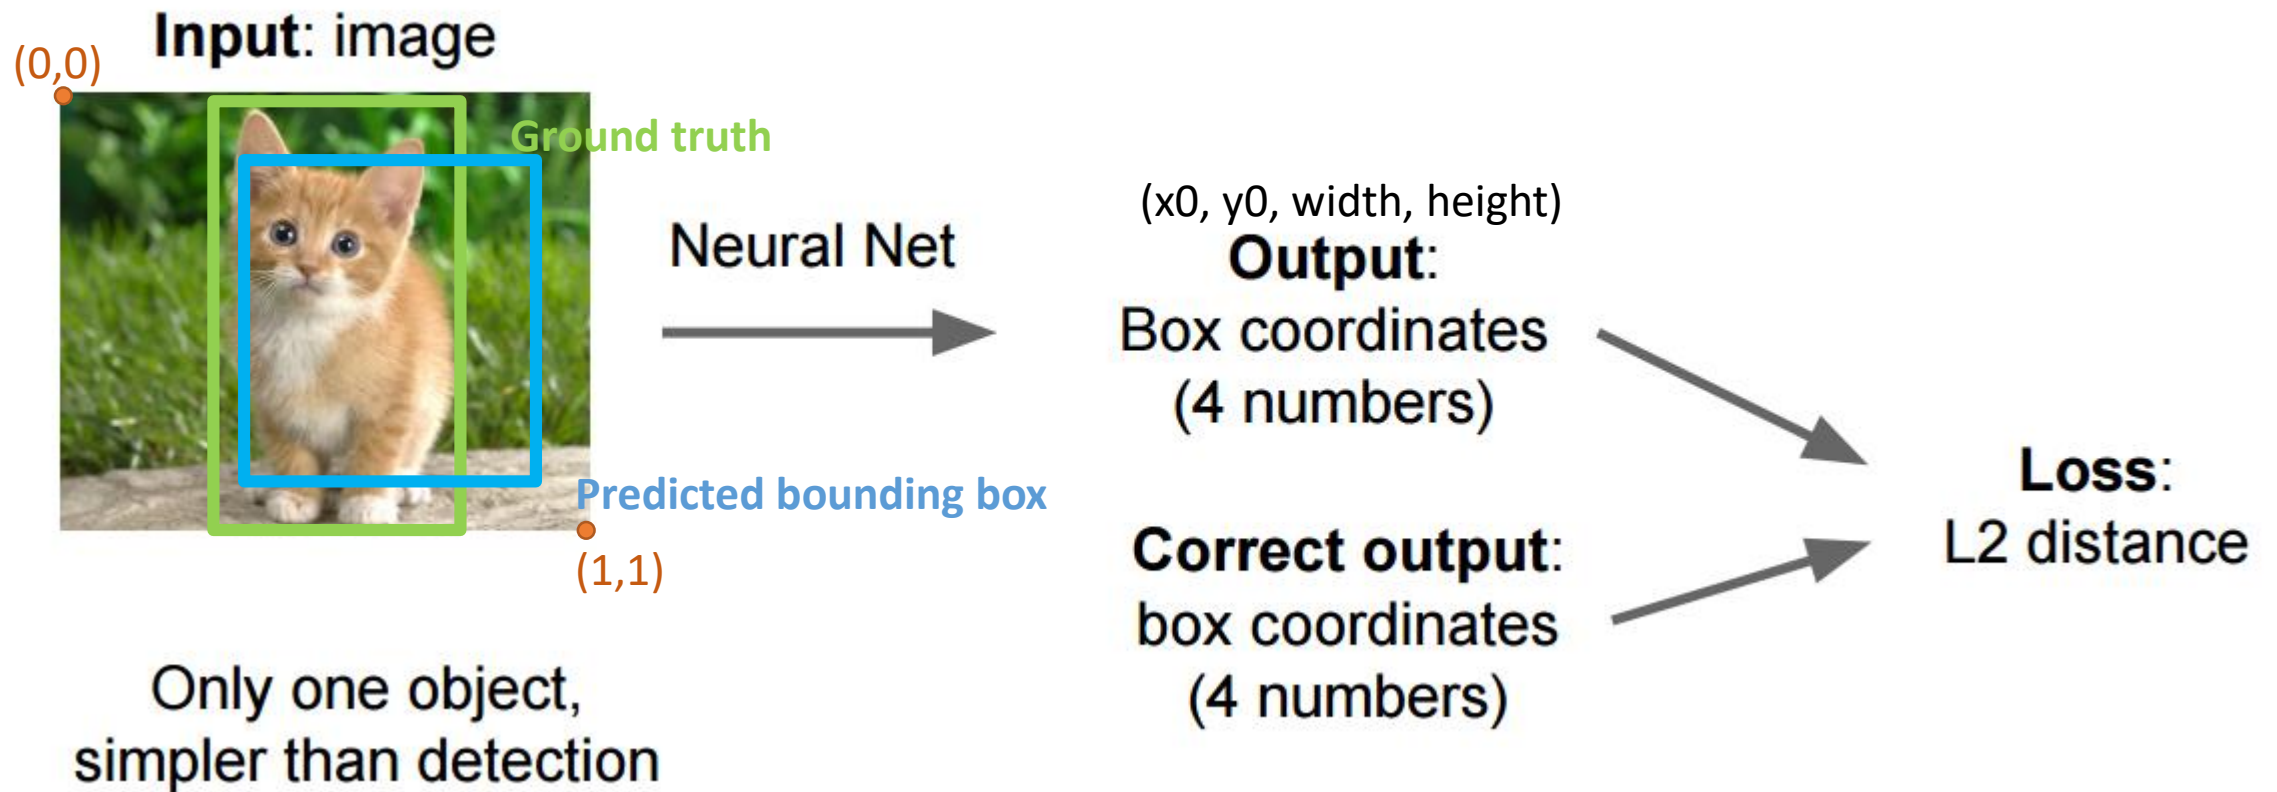
Instance Segmentation



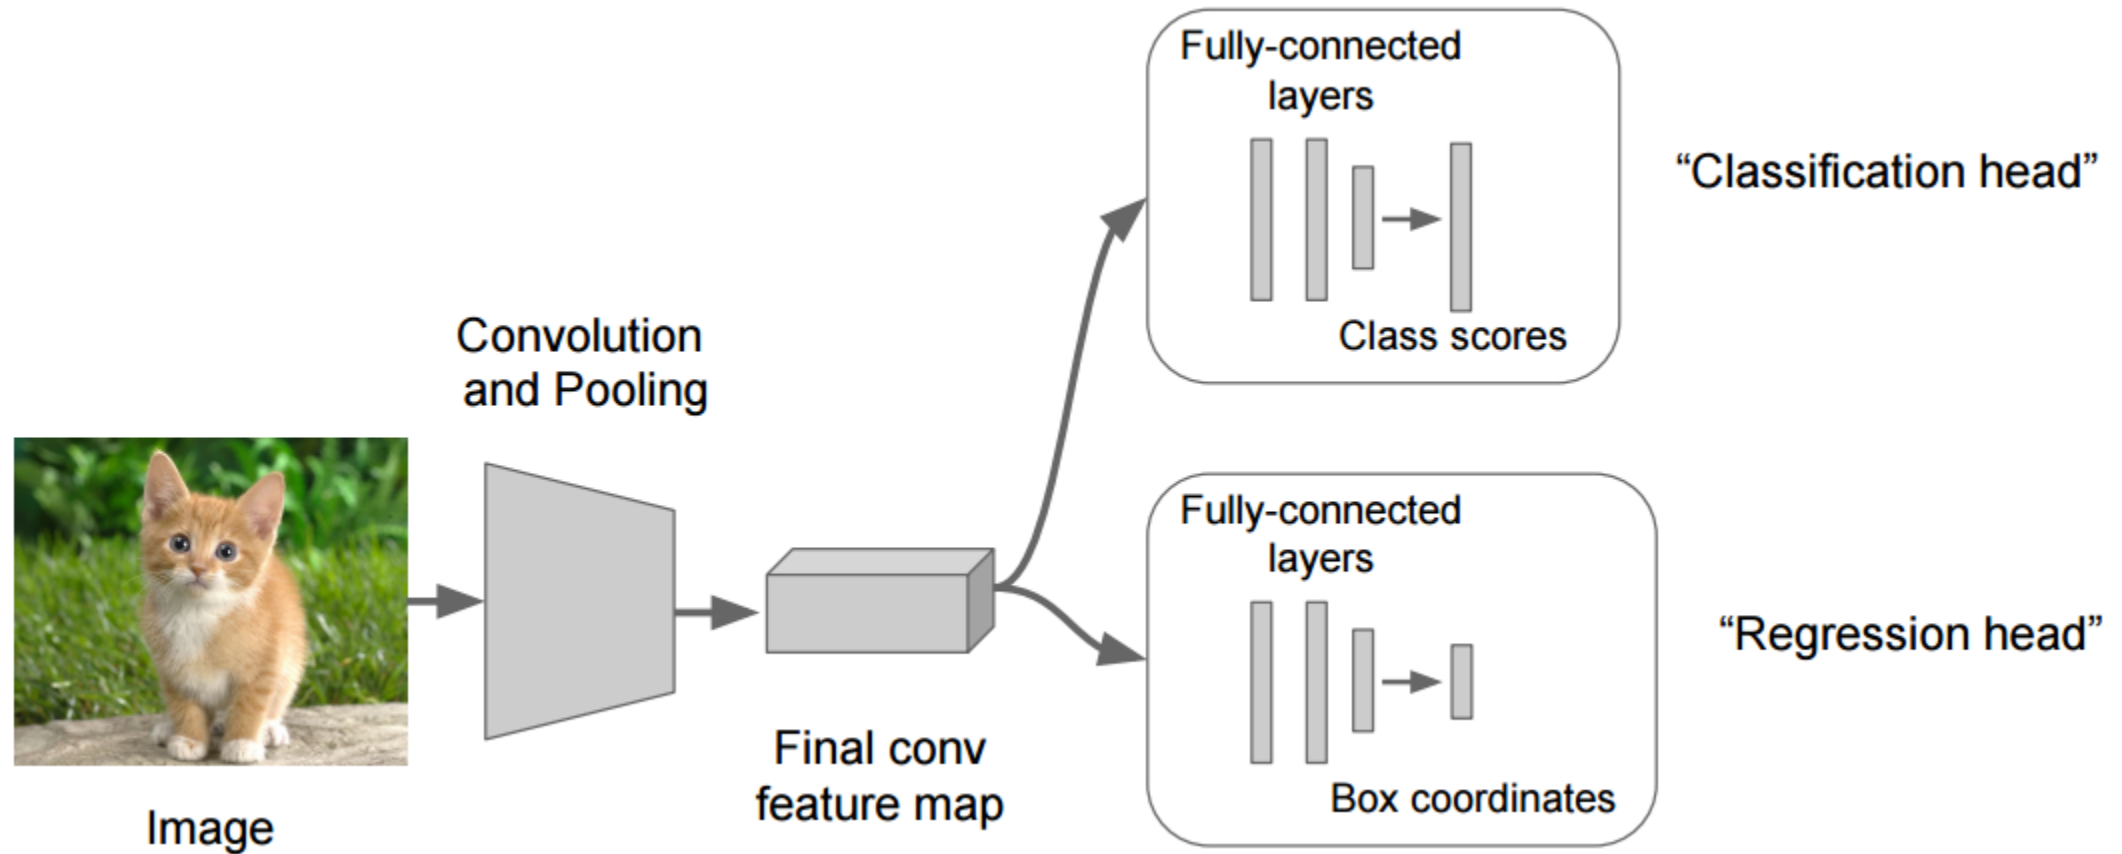
DOG, DOG, CAT

[This image is CC0 public domain](#)

Classification + Localization




Classification + Localization



Object Detection as Classification: Sliding Window



Apply a CNN to many different crops of the image, CNN classifies each crop as object or background

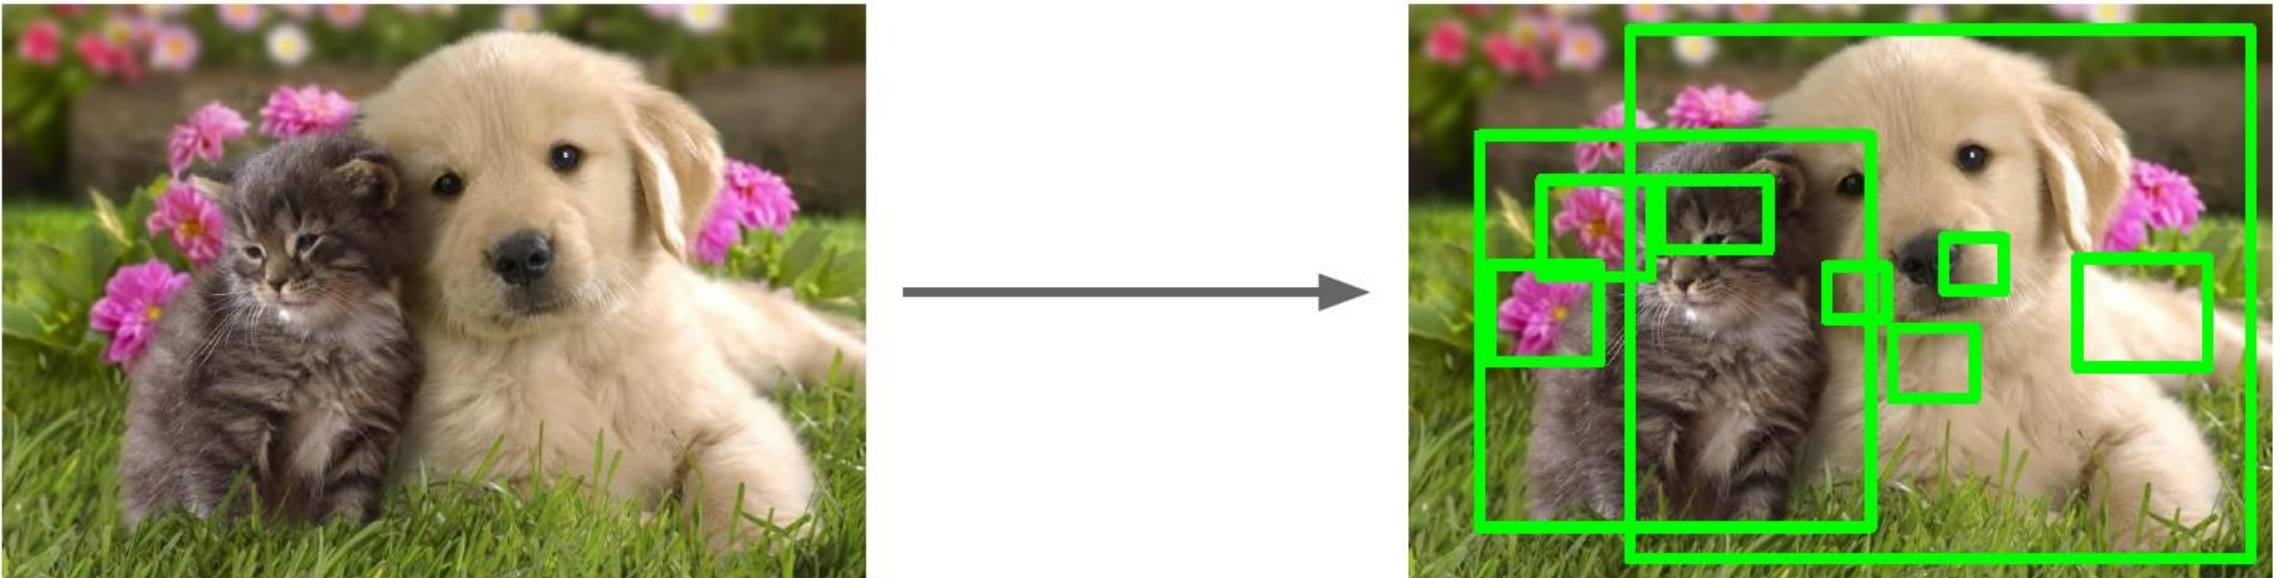


Dog? NO
Cat? YES
Background? NO

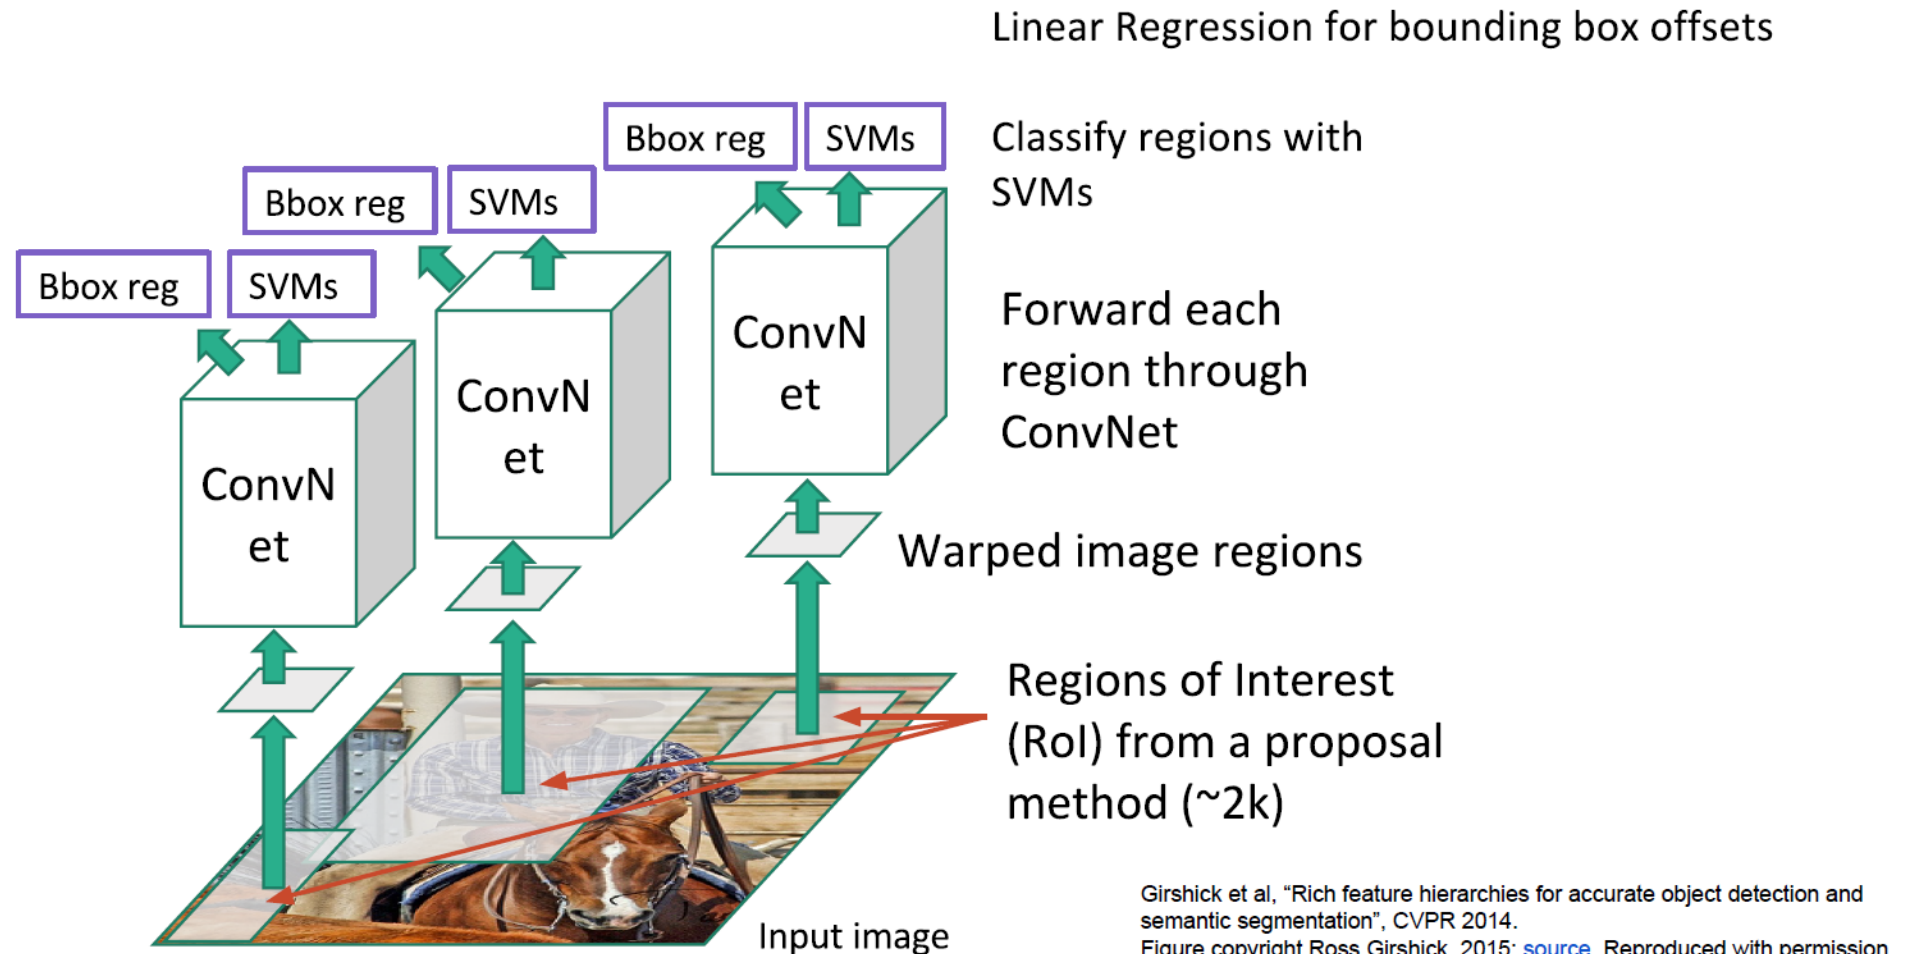
Problem: Need to apply CNN to huge number of locations and scales, very computationally expensive!

Region Proposals

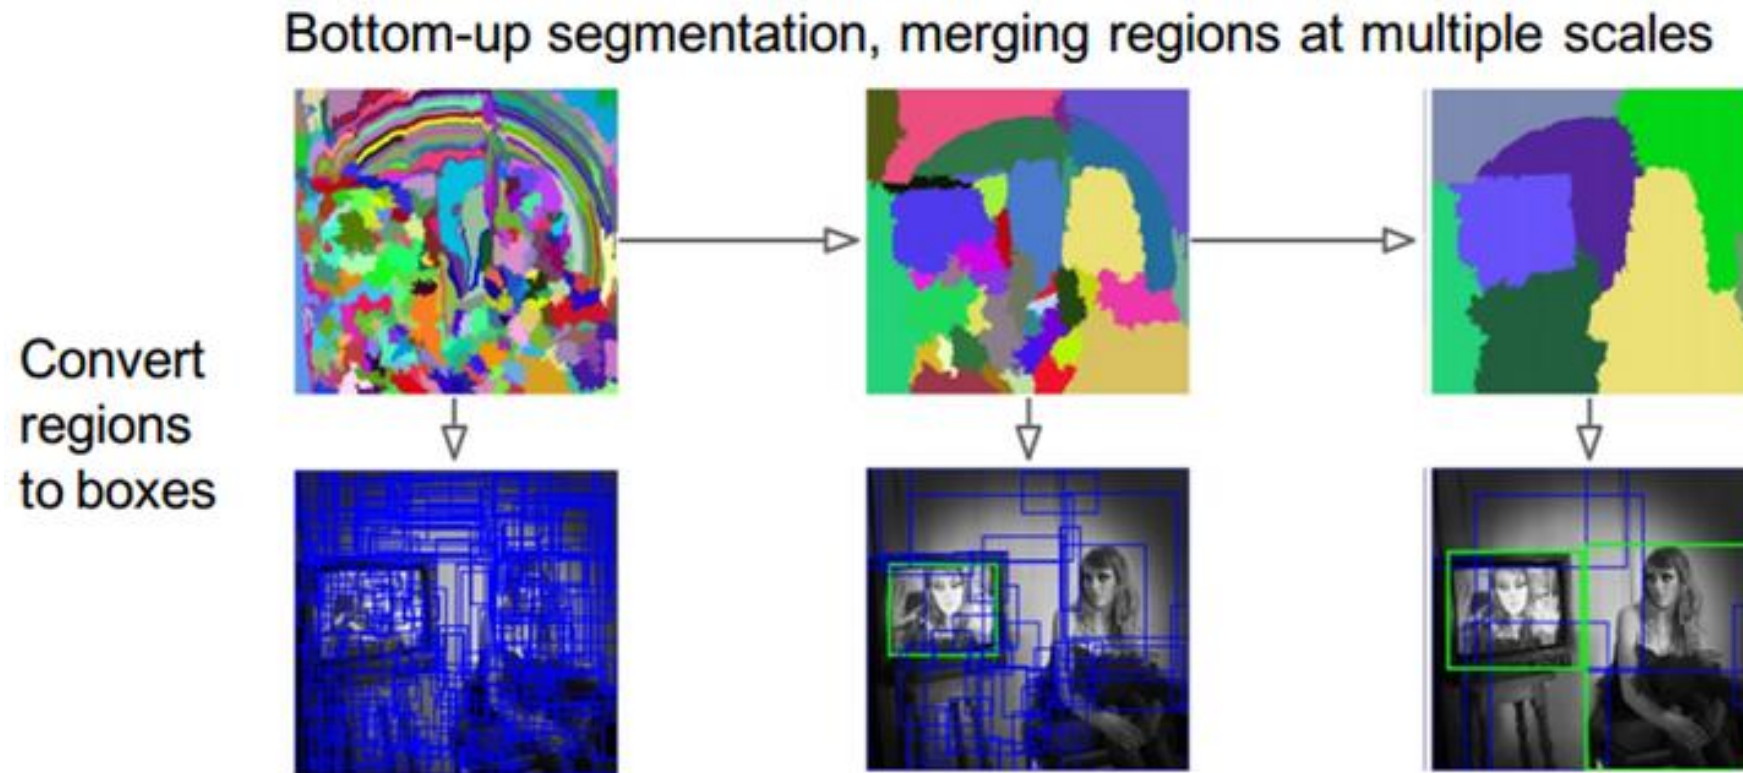
- Find image regions that are likely to contain objects
- To find regions, Selective Search algorithm gives 1K~2K region proposals per image



R-CNN: Regions with CNN features



Region Proposal – Selective Search



R-CNN Training

1. $M \leftarrow$ **Pre-train a ConvNet** for ImageNet classification dataset
2. $M' \leftarrow$ **Fine-tune** for object detection(softmax + log loss)
3. $R \leftarrow$ Selective Search for region proposals
4. $F \leftarrow$ **Cache feature vectors** to disk using M' in R
5. Train **post hoc** linear SVMs on F for object classification
6. Train **post hoc** linear bounding-box regressors on F

* **Post hoc** means the parameters are learned after the ConvNet is fixed

Bounding-Box Regression

$P^i = (P_x^i, P_y^i, P_w^i, P_h^i)$ Specifies the pixel coordinates of the center of proposal P^i 's
Bounding box together with P^i 's width and height

$G^i = (G_x^i, G_y^i, G_w^i, G_h^i)$ The ground-truth bounding box

$$\hat{G}_x = P_w d_x(P) + P_x \quad (1)$$

$$\hat{G}_y = P_h d_y(P) + P_y \quad (2)$$

$$\hat{G}_w = P_w \exp(d_w(P)) \quad (3)$$

$$\hat{G}_h = P_h \exp(d_h(P)). \quad (4)$$

Bounding-Box Regression

Each function $d_\star(P)$ (where \star is one of x, y, h, w) is modeled as a linear function of the pool_5 features of proposal P , denoted by $\phi_5(P)$. (The dependence of $\phi_5(P)$ on the image data is implicitly assumed.) Thus we have $d_\star(P) = \mathbf{w}_\star^T \phi_5(P)$, where \mathbf{w}_\star is a vector of learnable model parameters. We learn \mathbf{w}_\star by optimizing the regularized least squares objective (ridge regression):

$$\mathbf{w}_\star = \underset{\hat{\mathbf{w}}_\star}{\operatorname{argmin}} \sum_i^N (t_\star^i - \hat{\mathbf{w}}_\star^T \phi_5(P^i))^2 + \lambda \|\hat{\mathbf{w}}_\star\|^2. \quad (5)$$

$$\hat{G}_x = P_w d_x(P) + P_x \quad (1)$$

$$\hat{G}_y = P_h d_y(P) + P_y \quad (2)$$

$$\hat{G}_w = P_w \exp(d_w(P)) \quad (3)$$

$$\hat{G}_h = P_h \exp(d_h(P)). \quad (4)$$

The regression targets t_\star for the training pair (P, G) are defined as

$$t_x = (G_x - P_x)/P_w \quad (6)$$

$$t_y = (G_y - P_y)/P_h \quad (7)$$

$$t_w = \log(G_w/P_w) \quad (8)$$


$$t_h = \log(G_h/P_h). \quad (9)$$

Problems of R-CNN

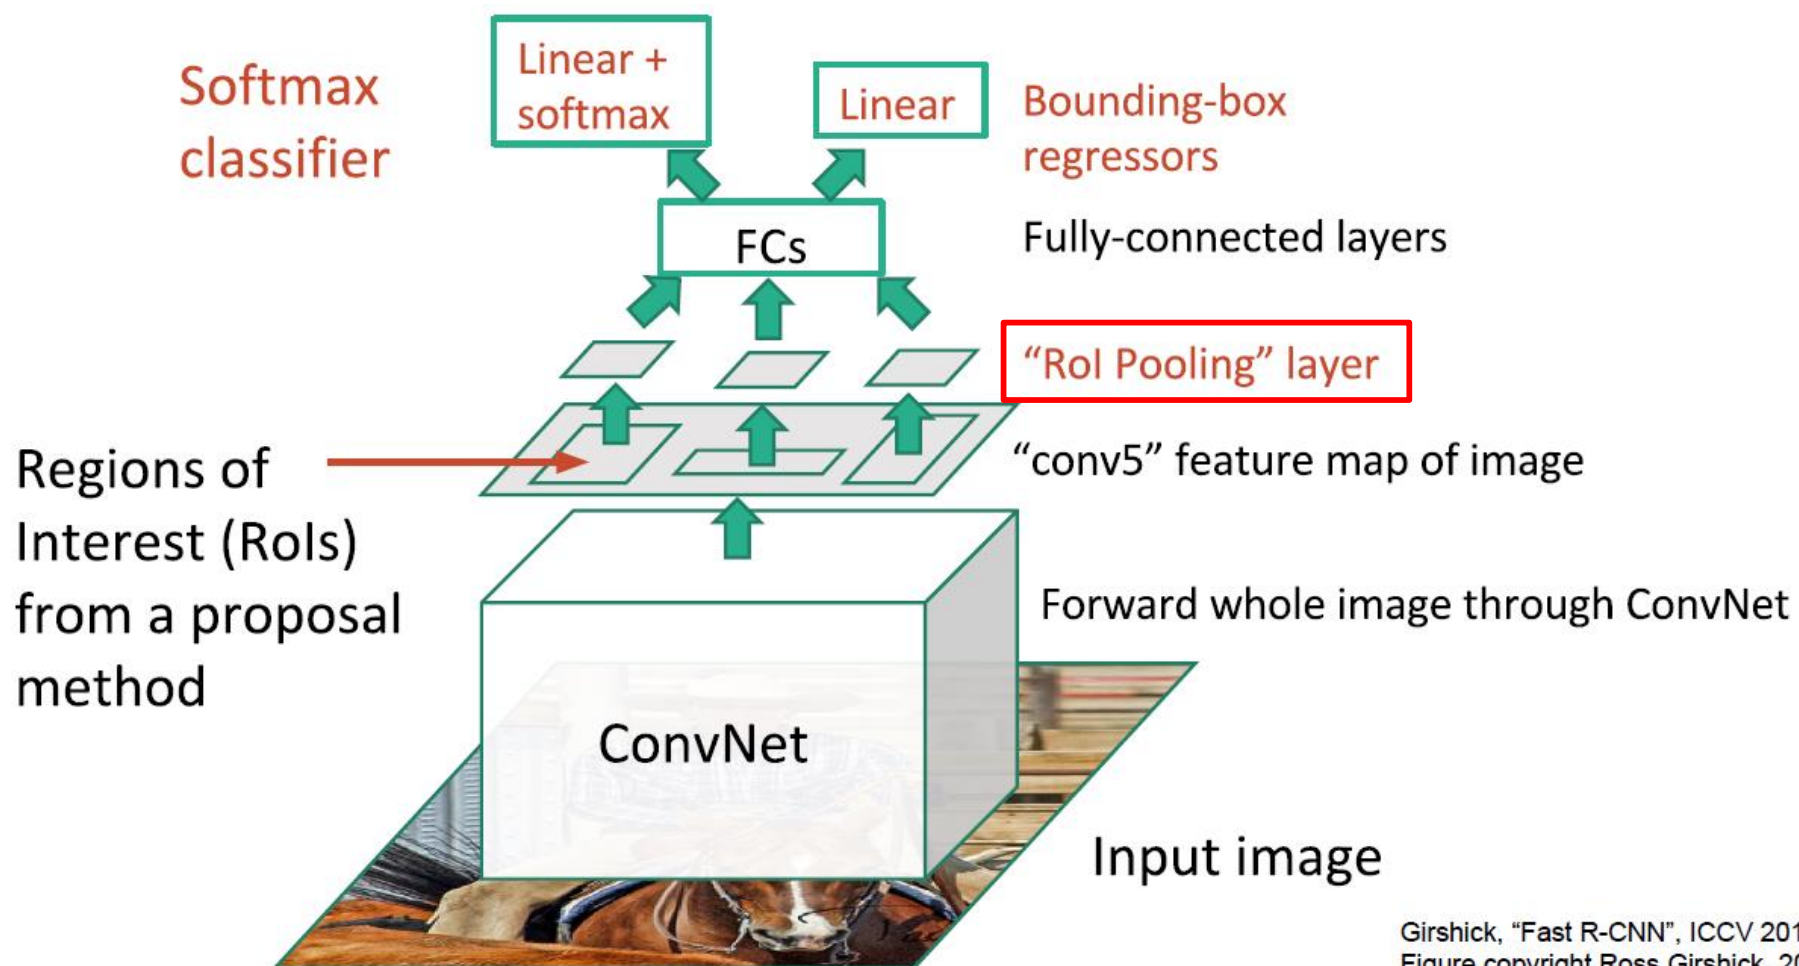
- **Slow at test-time**: need to run full forward path of CNN for each region proposal
 - 13s/image on a GPU
 - 53s/image on a CPU
- **SVM and regressors are post-hoc**: CNN features not updated in response to SVMs and regressors
- **Complex multistage training pipeline**
 - Fine-tune network with softmax classifier
 - Train post-hoc linear SVMs
 - Train post-hoc bounding-box regressions

Fast R-CNN

- Train the detector in a **single stage, end-to-end**
 - No caching features to disk
 - No post hoc training steps
- Train **all layers** of the network



Fast R-CNN

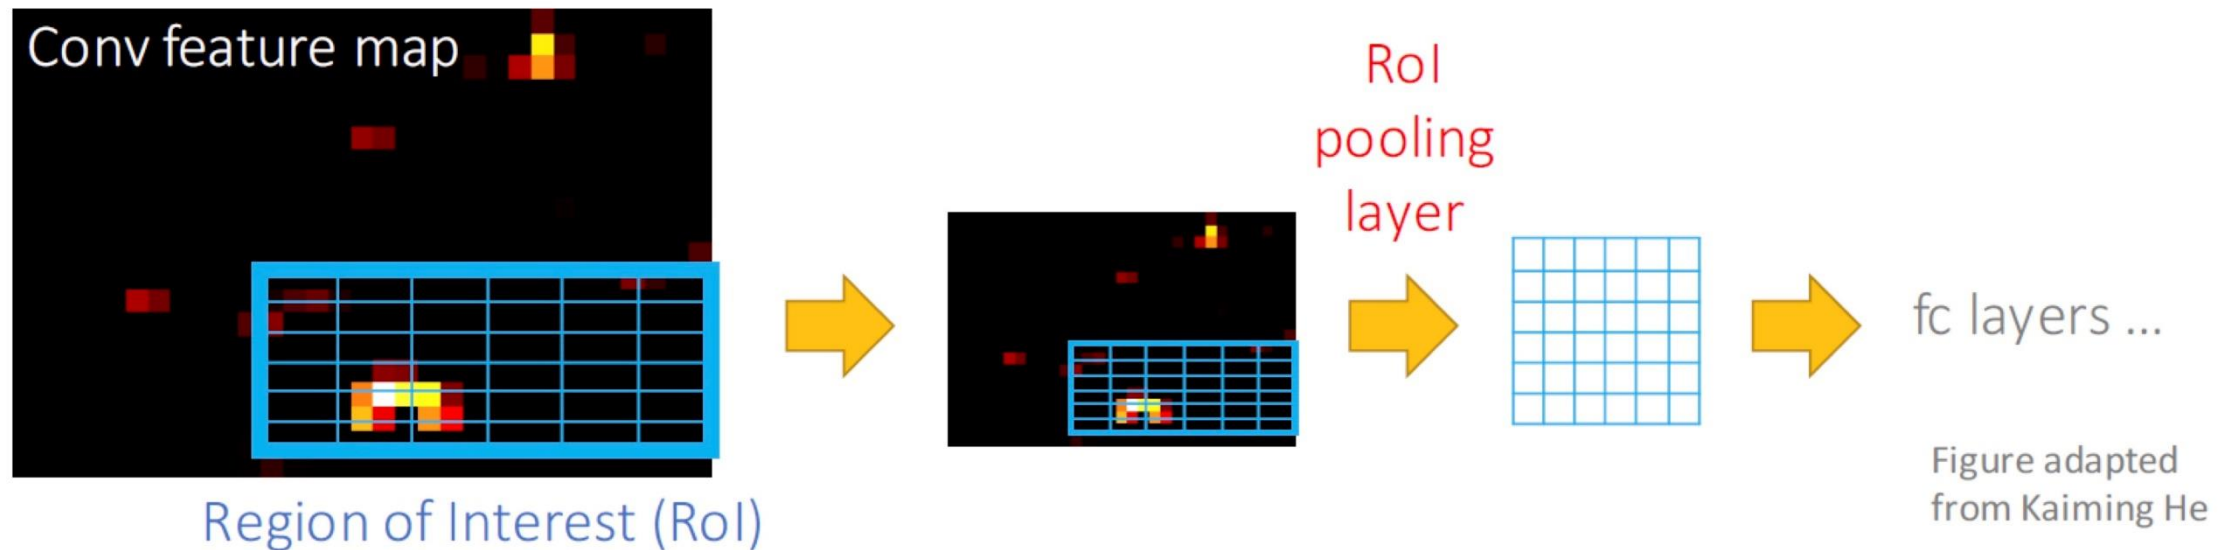


Girshick, "Fast R-CNN", ICCV 2015.

Figure copyright Ross Girshick, 2015; [source](#). Reproduced with permission.

ROI Pooling

- Its purpose is to perform max pooling on inputs of non uniform sizes to obtain fixed-size feature maps (e.g. 7×7 for VGG-net fc layer).
- The ROI pooling is done by:



ROI in Conv feature map : $21 \times 14 \rightarrow 3 \times 2$ max pooling with stride (3,2) \rightarrow output: 7×7

ROI in Conv feature map : $35 \times 42 \rightarrow 5 \times 6$ max pooling with stride (5,6) \rightarrow output: 7×7

Fast R-CNN Training

1. Take an input and a set of bounding boxes
2. Generate convolutional feature maps
3. For each bbox, get a fixed-length feature vector from ROI pooling layer
4. Outputs have two information
 1. K+1 class labels
 2. Bounding box locations

- Loss function

True box coordinates
Predicted box coordinates

True class score
Predicted class score

Log loss
Smooth L1 loss

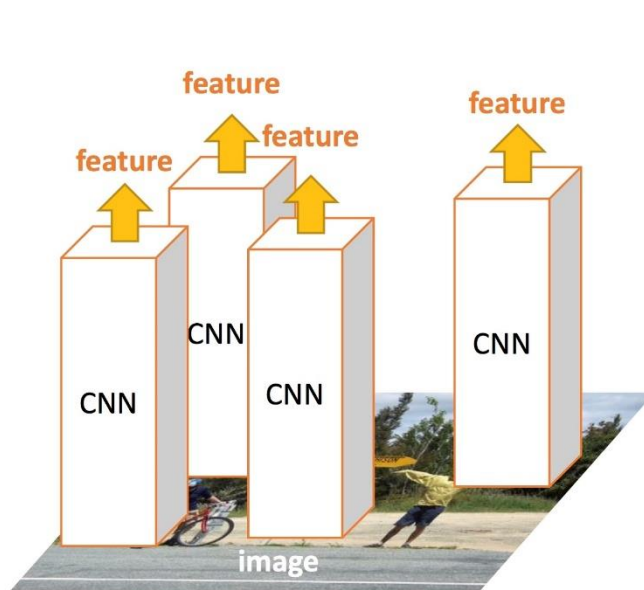
$$L(p, u, t^u, v) = L_{\text{cls}}(p, u) + \lambda[u \geq 1]L_{\text{loc}}(t^u, v), \quad (1)$$

in which

$$L_{\text{loc}}(t^u, v) = \sum_{i \in \{x, y, w, h\}} \text{smooth}_{L_1}(t_i^u - v_i), \quad (2)$$

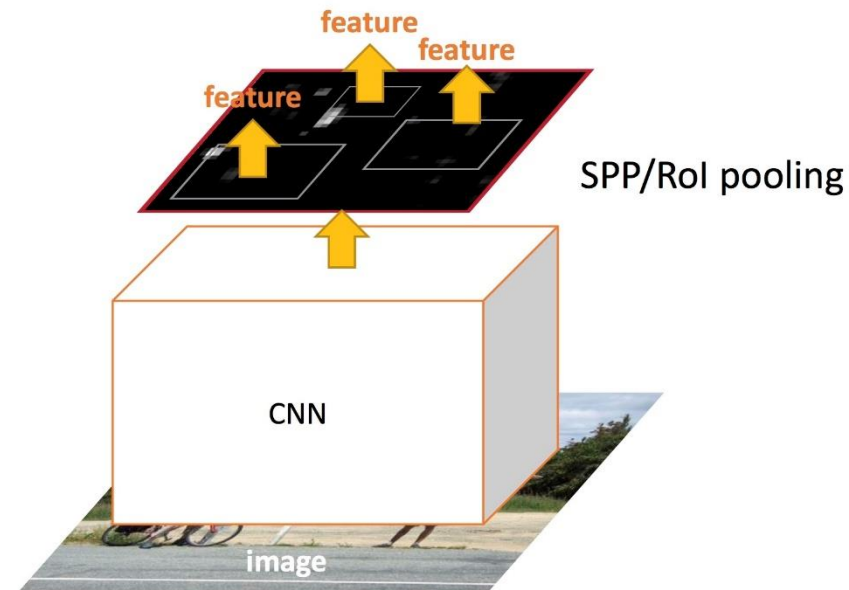
$$\text{smooth}_{L_1}(x) = \begin{cases} 0.5x^2 & \text{if } |x| < 1 \\ |x| - 0.5 & \text{otherwise,} \end{cases} \quad (3)$$

R-CNN vs Fast R-CNN



R-CNN

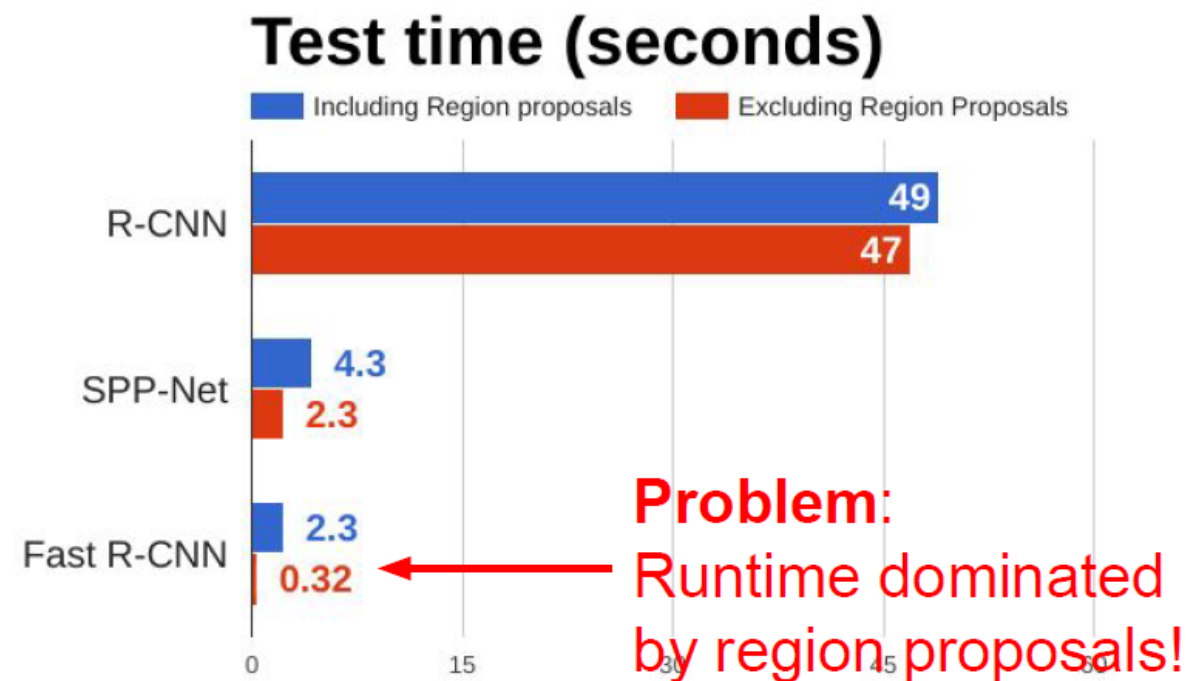
- Extract image regions
- 1 CNN per region (2000 CNNs)
- Classify region-based features



SPP-net & Fast R-CNN (the same forward pipeline)

- 1 CNN on the entire image
- Extract features from feature map regions
- Classify region-based features

R-CNN vs Fast R-CNN

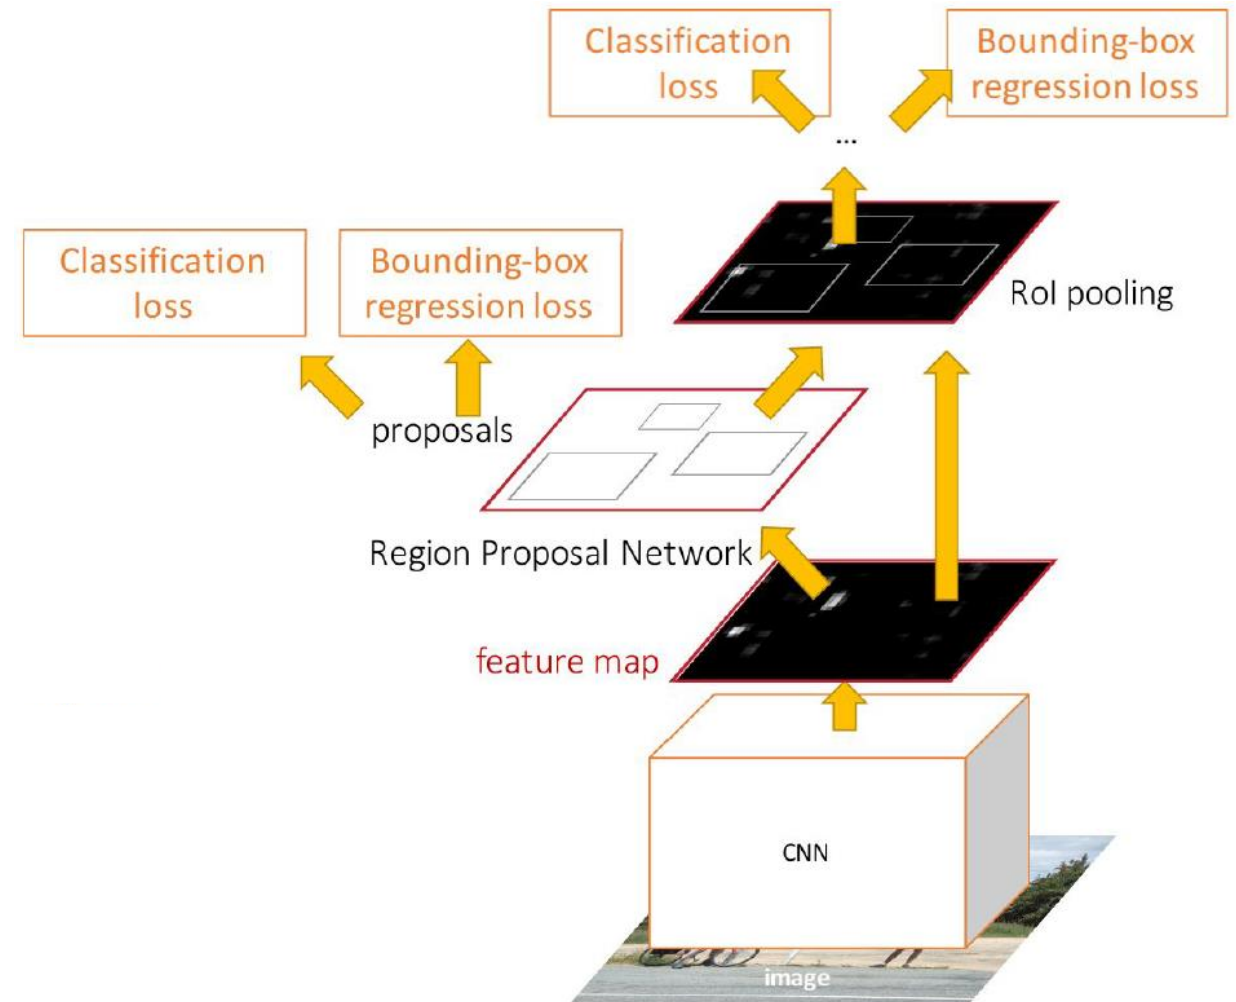


Problems of Fast R-CNN

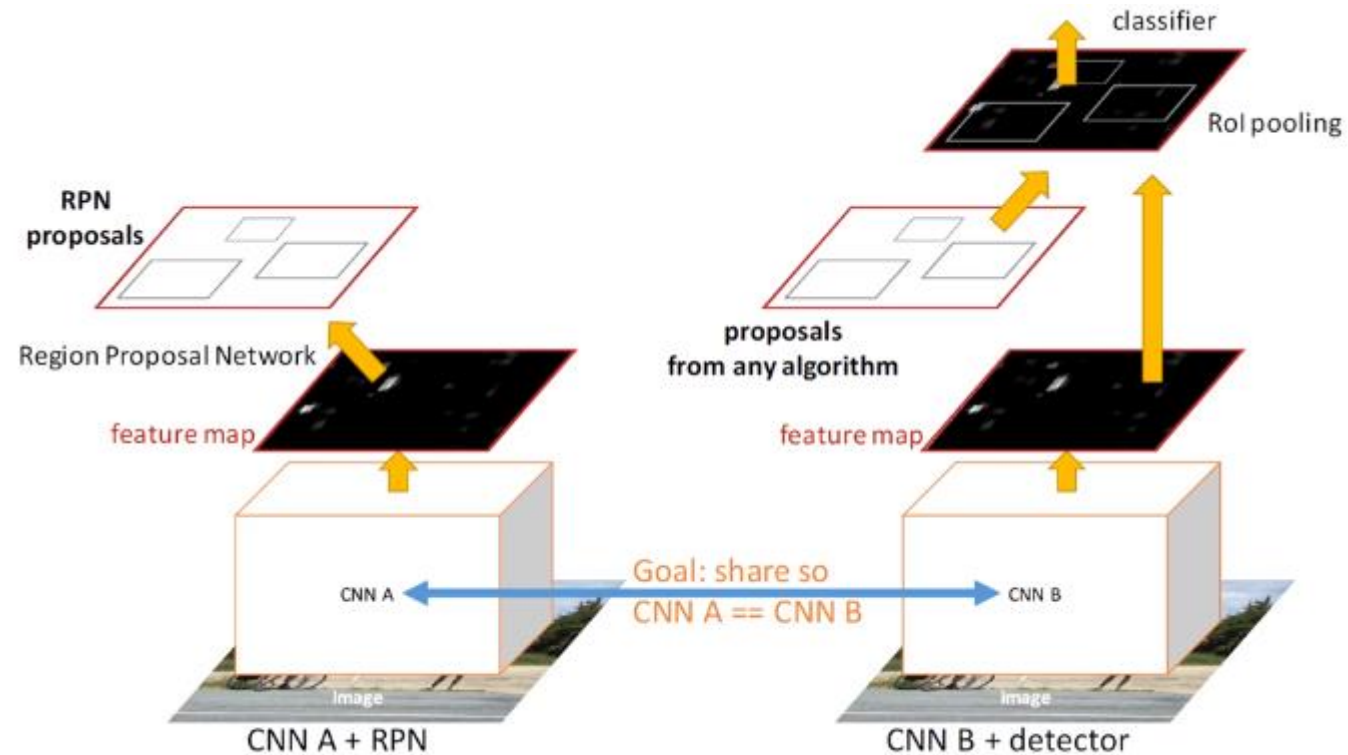
- Out-of-network region proposals are the test-time computational bottleneck
- Object detection networks are fast but,
 - Region proposal
 - Selective Search [Uijlings et al. ICCV 2011]: 2s per image
 - EdgeBoxes [Zitnick & Dollar. ECCV 2014]: 0.2s per image

Faster R-CNN (RPN + Fast R-CNN)

- Insert a Region Proposal Network (RPN) after the last convolutional layer -> Using GPU!
- RPN trained to produce region proposals directly; no need for external region proposals
- After RPN, use ROI Pooling and an upstream classifier and bbox regressor just like Fast R-CNN

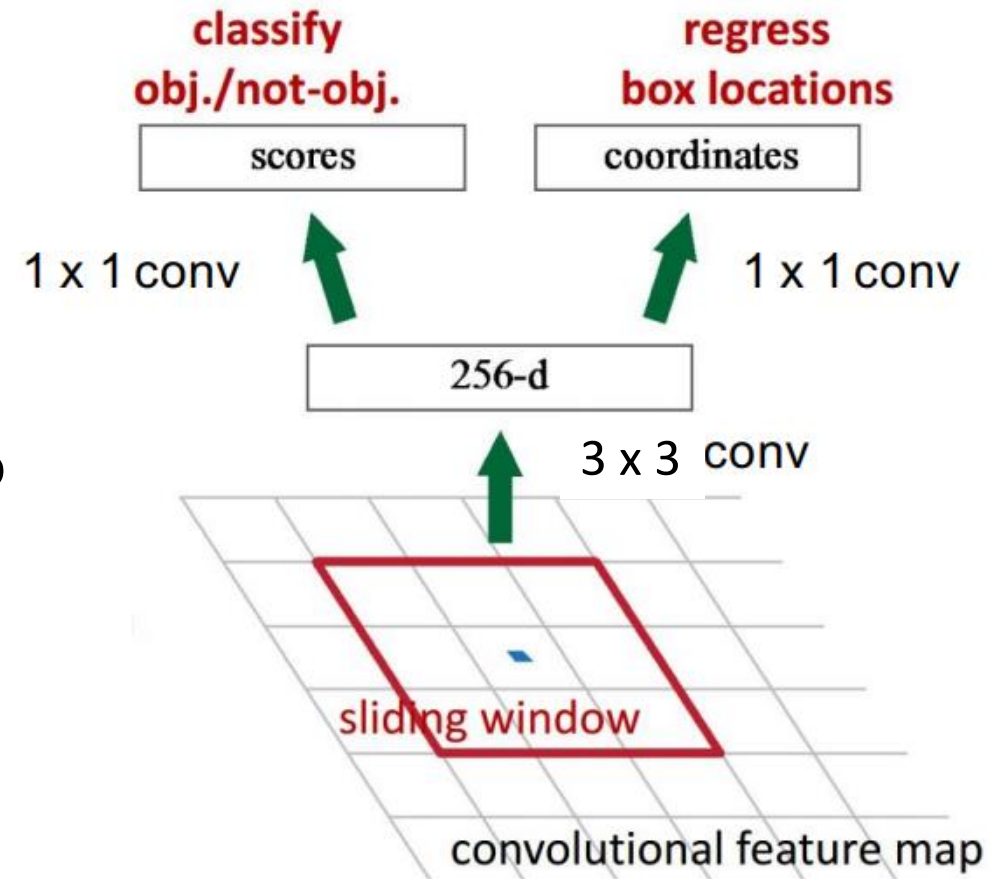


Training Goal: Share Feature Map



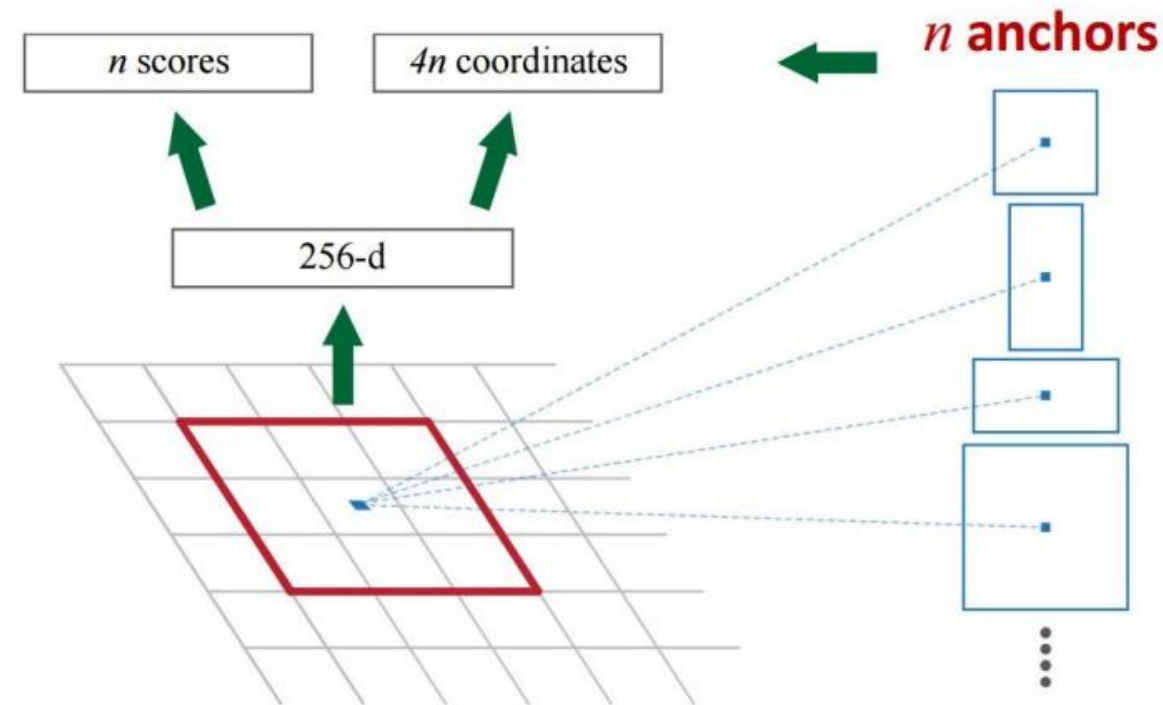
Faster R-CNN: Region Proposal Network

- Slide a small window on the feature map
- Build a small network for
 - Classifying object or not-object, and
 - Regressing bbox locations
- Position of the sliding window provides localization information with reference to the image
- Box regression provides finer localization information with reference to this sliding window



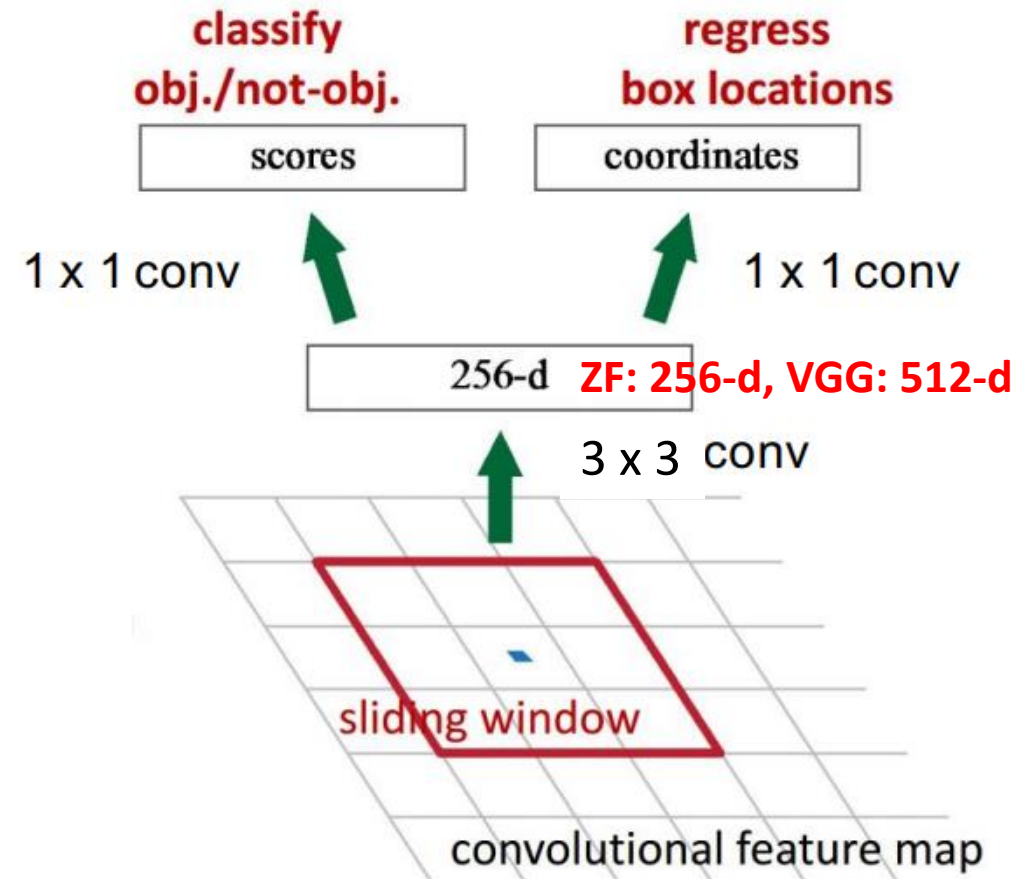
Faster R-CNN: Region Proposal Network

- Use N anchor boxes at each location
- Anchors are translation invariant: use the same ones at every location
- Regression gives offsets from anchor boxes
- Classification gives the probability that each (regressed) anchor shows an object



Faster R-CNN: Region Proposal Network

- Intermediate layer – 256(or 512) 3x3 filter, stride 1, padding 1
- Cls layer – 18(9x2) 1x1 filter, stride 1, padding 0
- Reg layer – 36(9x4) 1x1 filter, stride 1, padding 0



Anchors

- Anchors: pre-defined reference boxes
- Multi-scale/size anchors:
 - Multiple anchors are used at each position:
 - 3 scale(128x128, 256x256, 512x512), 3 aspect ratios (2:1, 1:1, 1:2) yield 9 anchors
 - Each anchor has its own prediction function

Positive/Negative Samples

- An anchor is **labeled as positive** if
 - The anchor is the one with **highest IoU** overlap with a ground-truth box
 - The anchor has an IoU overlap with a ground-truth box **higher than 0.7**
- **Negative labels** are assigned to anchors with IoU **lower than 0.3** for all ground-truth boxes
- 50%/50% ratio of positive/negative anchors in a minibatch

RPN Loss Function

i : index of an anchor in a mini-batch

$$L(\{p_i\}, \{t_i\}) = \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) + \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*)$$

p_i : predicted probability of anchor i being an object

t_i : vector representing the 4 parameterized coordinates of the predicted bounding box

p_i^* : 1 if the anchor is positive, and 0 if the anchor is negative

t_i^* : ground-truth box associated with a positive anchor

L_{cls} : log loss over two classes (object vs. not object)

L_{reg} : smooth L1 loss

N_{cls} : number of anchors in mini-batch (~256)

N_{reg} : number of anchor locations (~2400)

λ : default is 10, so that both terms are roughly equally balanced

4-Step RPN Training

- M0 is an ImageNet pre-trained network

1. `train_rpn(M0) -> M1`

2. `generate_proposals(M1) -> P1`

3. `train_fast_rcnn(M0, P1) -> M2`

4. `train_rpn_frozen_conv(M2) -> M3`

5. `generate_proposals(M3) -> P2`

6. `train_fast_rcnn_frozen_conv(M3, P2) -> M4`

7. `return add_rpn_layers(M4, M3.RPN)`

Train an RPN initialized from M0, get M1


Generate training proposals P1 using RPN M1

Train Fast R-CNN M2 on P1 initialized from M0

Train RPN M3 from M2 without changing conv layers

Conv layers are shared with RPN M3

Add M3's RPN layers to Fast R-CNN M4



Results

	R-CNN	Fast R-CNN	Faster R-CNN
Test time per image (with proposals)	50 sec	2 sec	0.2 sec
Speedup	1x	25x	250x
mAP (VOC 2007)	66.0	66.9	66.9

Table 8: Detection results of Faster R-CNN on PASCAL VOC 2007 test set using **different settings of anchors**. The network is VGG-16. The training data is VOC 2007 trainval. The default setting of using 3 scales and 3 aspect ratios (69.9%) is the same as that in Table 3.

settings	anchor scales	aspect ratios	mAP (%)
1 scale, 1 ratio	128^2	1:1	65.8
	256^2	1:1	66.7
1 scale, 3 ratios	128^2	{2:1, 1:1, 1:2}	68.8
	256^2	{2:1, 1:1, 1:2}	67.9
3 scales, 1 ratio	{ $128^2, 256^2, 512^2$ }	1:1	69.8
3 scales, 3 ratios	{ $128^2, 256^2, 512^2$ }	{2:1, 1:1, 1:2}	69.9

Table 9: Detection results of Faster R-CNN on PASCAL VOC 2007 test set using **different values of λ** in Equation (1). The network is VGG-16. The training data is VOC 2007 trainval. The default setting of using $\lambda = 10$ (69.9%) is the same as that in Table 3.

λ	0.1	1	10	100
mAP (%)	67.2	68.9	69.9	69.1

Problems of Faster R-CNN

- ROI Pooling has some quantization operations
- These operations introduce misalignments between the ROI and the extracted features
- While this may not impact classification, it can make a negative effect on predicting bounding-box

=> **MASK R-CNN**

Check the code

- https://github.com/smallcorgi/Faster-RCNN_TF

References

- Object detection
 - R-CNN: [“Rich feature hierarchies for accurate object detection and semantic segmentation”](#) (2013)
 - SPPNet: [“Spatial Pyramid Pooling in Deep Convolutional Networks for Visual Recognition”](#) (2014)
 - Fast R-CNN: [“Fast R-CNN”](#) (2015)
 - Faster R-CNN: [“Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks”](#) (2015)
 - SSD: [“SSD: Single Shot MultiBox Detector”](#) (2015)
 - Mask R-CNN: [Mask R-CNN](#) (2017)
- Others
 - Selective Search: [“Selective Search for Object Recognition”](#) (2012)
- Github
 - R-CNN
 - (Caffe + MATLAB): <https://github.com/rbgirshick/rcnn>
 - Fast R-CNN
 - (Caffe + MATLAB): <https://github.com/rbgirshick/fast-rcnn>
 - Faster R-CNN
 - (tensorflow) https://github.com/smallcorgi/Faster-RCNN_TF